

# Principal Component Analysis All in One

Version 0.4.9

Zilong Li\*

December 26, 2024

## Contents

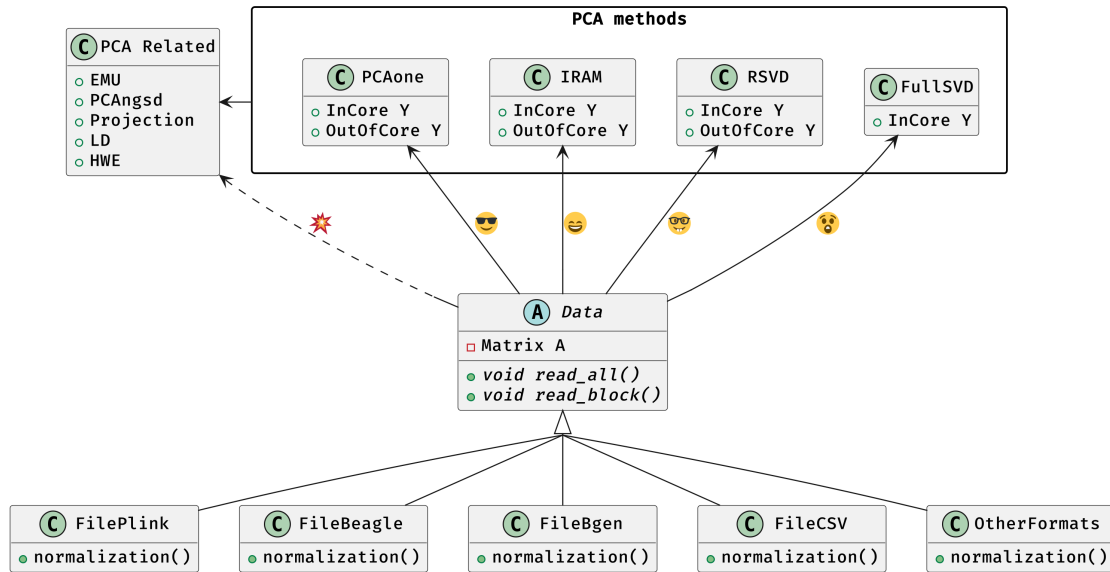
<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Features</b>	<b>2</b>
<b>3</b>	<b>Cite the work</b>	<b>3</b>
<b>4</b>	<b>Quick start</b>	<b>3</b>
<b>5</b>	<b>Installation</b>	<b>4</b>
5.1	Download compiled binary . . . . .	4
5.2	Via Conda . . . . .	4
5.3	Build from source . . . . .	4
<b>6</b>	<b>Documentation</b>	<b>6</b>
6.1	Options . . . . .	6
6.2	Which SVD method to use . . . . .	8
6.3	Input formats . . . . .	8
6.4	Output formats . . . . .	8
6.5	Memory-efficient modes . . . . .	9
6.6	Data Normalization . . . . .	9
6.7	Projection . . . . .	10
6.8	HWE accounting for population structure . . . . .	10
6.9	Ancestry-Adjusted LD matrix . . . . .	10
6.10	Report LD statistics . . . . .	10
6.11	Pruning based on Ancestry-Adjusted LD . . . . .	11
6.12	Clumping based on Ancestry-Adjusted LD . . . . .	11
<b>7</b>	<b>More tutorials</b>	<b>11</b>
7.1	Genotype data (PLINK) . . . . .	11
7.2	Genotype dosage (BGEN) . . . . .	11
7.3	Single cell RNA-seq data (CSV) . . . . .	12
<b>8</b>	<b>Acknowledgements</b>	<b>12</b>

---

\*zilong.dk@gmail.com

# 1 Introduction

PCAone is a fast and memory efficient PCA tool implemented in C++ aiming at providing comprehensive features and algorithms for different scenarios. PCAone implements 3 fast PCA algorithms for finding the top eigenvectors of large datasets, which are **Implicitly Restarted Arnoldi Method** (IRAM, `-svd 0`), **single pass Randomized SVD** but with power iteration scheme (RSVD, `-svd 1`, Algorithm1 in paper) and **our own RSVD with window based power iteration scheme** (PCAone, `-svd 2`, Algorithm2 in paper). All have both in-core and out-of-core implementation. Additionally, full SVD (`-svd 3`) is supported via in-core mode only. There is also an **R** package here. PCAone supports multiple different input formats, such as **PLINK**, **BGEN**, **Beagle** genetic data formats and a general comma separated CSV format for other data, such as scRNAs and bulk RNAs. For genetics data, PCAone also implements **EMU** and **PCAngsd** algorithm for data with missingness and uncertainty. The PDF manual can be downloaded [here](#).



## 2 Features

See [change log](#) here.

- Has both Implicitly Restarted Arnoldi Method (IRAM) and Randomized SVD (RSVD) with **out-of-core** implementation.
- Implements our new fast window based Randomized SVD algorithm for tera-scale dataset.
- Quite fast with multi-threading support using high performance library **MKL** or **OpenBLAS** as backend.
- Supports the **PLINK**, **BGEN**, **Beagle** genetic data formats.
- Supports a general comma separated CSV format for single cell RNA-seq or bulk RNA-seq data compressed by **zstd**.
- Supports **EMU** algorithm for scenario with large proportion of missingness.
- Supports **PCAngsd** algorithm for low coverage sequencing scenario with genotype likelihood as input.
- Novel **LD** pruning and clumping method that accounts for population structure in the data.

- Projection support for data with missingness.
- HWE test taking population structure into account.

### 3 Cite the work

- If you use PCAone, please first cite our paper on genome reseach [Fast and accurate out-of-core PCA framework for large scale biobank data](#).
- If using the EMU algorithm, please also cite [Large-scale inference of population structure in presence of missingness using PCA](#).
- If using the PCAngsd algorithm, please also cite [Inferring Population Structure and Admixture Proportions in Low-Depth NGS Data](#).
- If using the ancestry ajusted LD statistics for pruning and clumping, please also cite [Measuring linkage disequilibrium and improvement of pruning and clumping in structured populations](#).

### 4 Quick start

```
pkg=https://github.com/Zilong-Li/PCAone/releases/latest/download/PCAone-avx2-Linux.zip
wget $pkg
unzip -o PCAone-avx2-Linux.zip
wget http://popgen.dk/zilong/datahub/pca/example.tar.gz
tar -xzf example.tar.gz && rm -f example.tar.gz
# in default calculating top 10 PCs with in-memory mode
./PCAone -b example/plink
R -s -e 'd=read.table("pcaone.eigvecs2", h=F);
plot(d[,1:2+2], col=factor(d[,1]), xlab="PC1", ylab="PC2");
legend("topright", legend=unique(d[,1]), col=1:4, pch = 21, cex=1.2);'
```

You will find these files in current directory.

```
.
PCAone          # program
Rplots.pdf      # pca plot
example         # folder of example data
pcaone.eigvals  # eigenvalues
pcaone.eigvecs  # eigenvectors, the PCs you need to plot
pcaone.eigvecs2 # eigenvectors with header line
pcaone.log      # log file
```

## 5 Installation

There are 3 ways to install PCAone.

### 5.1 Download compiled binary

There are compiled binaries provided for both Linux and Mac platform. Check [the releases page](#) to download one. For most modern CPUs and Linux systems, download the one named with `avx2`.

```
pkg=https://github.com/Zilong-Li/PCAone/releases/latest/download/PCAone-avx2-Linux.zip
wget $pkg || curl -LO $pkg
unzip -o PCAone-avx2-Linux.zip
```

If the server is too old to support `avx2` instruction, download the alternative version.

```
pkg=https://github.com/Zilong-Li/PCAone/releases/latest/download/PCAone-x64-Linux.zip
wget $pkg || curl -LO $pkg
unzip -o PCAone-x64-Linux.zip
```

### 5.2 Via Conda

PCAone is also available from [bioconda](#).

```
conda config --add channels bioconda
conda install pcaone
PCAone --help
```

### 5.3 Build from source

PCAone can be running on a normal computer/laptop with `x86-64` instruction set architecture. PCAone has been tested on both Linux and MacOS system. To build PCAone from the source code, the following dependencies are required:

- GCC/Clang compiler with C++17 support
- GNU make
- zlib

We **recommend** building the software from source with MKL as backend to maximize the performance. For MacOS users, we recommend using `llvm` by `brew install llvm` instead of the default `clang` shipped with MacOS. Check out the [mac workflow](#).

#### 5.3.1 With MKL or OpenBLAS as backend

Build PCAone dynamically with MKL can maximize the performance since the faster threading layer `libiomp5` will be linked at runtime. One can obtain the MKL by one of the following option:

- install `mkl` by conda

```
conda install -c conda-forge -c anaconda -y mkl mkl-include intel-openmp
git clone https://github.com/Zilong-Li/PCAone.git
cd PCAone
# if mkl is installed by conda then use ${CONDA_PREFIX} as mklroot
make -j4 MKLROOT=${CONDA_PREFIX}
./PCAone -h
```

- download `mkl` from [the website](#)

After having `mkl` installed, find the `mkl` root path and replace the path below with your own.

```
# if libiomp5 is not in the mklroot path, please link it to $MKLROOT/lib folder
make -j4 MKLROOT=/path/to/mklroot
```

Alternatively, for advanced user, modify variables directly in **Makefile** and run **make** to use MKL or OpenBlas as backend.

### 5.3.2 Without MKL or OpenBLAS dependency

If you don't want any optimized math library as backend, just run:

```
git clone https://github.com/Zilong-Li/PCAone.git
cd PCAone
make -j4
./PCAone -h
```

If this doesn't work because the server is too outdated, run **make clean && make -j4 AVX=0** instead.

## 6 Documentation

### 6.1 Options

Run `PCAone --groff > pcaone.1 && man ./pcaone.1` or `PCAone --help` to read the manual. Here are common options.

```
General options:
-h, --help                print all options including hidden advanced options
-m, --memory arg (=0)    RAM usage in GB unit for out-of-core mode. default is in-core
    ↪ mode
-n, --threads arg (=12)  the number of threads to be used
-v, --verbose arg (=1)   verbosity level for logs. any level x includes messages for all
    ↪ levels (1...x). Options are
                        0: silent, no messages on screen;
                        1: concise messages to screen;
                        2: more verbose information;
                        3: enable debug information.

PCA algorithms:
-d, --svd arg (=2)       SVD method to be applied. default 2 is recommended for big data.
    ↪ Options are
                        0: the Implicitly Restarted Arnoldi Method (IRAM);
                        1: the Yu's single-pass Randomized SVD with power iterations;
                        2: the accurate window-based Randomized SVD method (PCAone);
                        3: the full Singular Value Decomposition.
-k, --pc arg (=10)       top k principal components (PCs) to be calculated
-C, --scale arg (=0)     do scaling for input file. Options are
                        0: do just centering;
                        1: do log transformation eg. log(x+0.01) for RNA-seq data;
                        2: do count per median log transformation (CPMED) for scRNAs.
-p, --maxp arg (=40)     maximum number of power iterations for RSVD algorithm.
-S, --no-shuffle         do not shuffle columns of data for --svd 2 (if not locally
    ↪ correlated).
--emu                   use EMU algorithm for genotype input with missingness.
--pcangsd               use PCAngsd algorithm for genotype likelihood input.

Input options:
-b, --bfile arg          prefix of PLINK .bed/.bim/.fam files.
-B, --binary arg        path of binary file.
-c, --csv arg            path of comma seperated CSV file compressed by zstd.
-g, --bgen arg          path of BGEN file compressed by gzip/zstd.
-G, --beagle arg        path of BEAGLE file compressed by gzip.
-f, --match-bim arg     the .mbim file to be matched, where the 7th column is allele
    ↪ frequency.
--USV arg               prefix of PCAone .eigvecs/.eigvals/.loadings/.mbim.

Output options:
-o, --out arg (=pcaone)  prefix of output files. default [pcaone].
-V, --printv            output the right eigenvectors with suffix .loadings.
-D, --ld               output a binary matrix for downstream LD related analysis.
-R, --print-r2          print LD R2 to *.ld.gz file for pairwise SNPs within a window
    ↪ controlled by --ld-bp.

Misc options:
--maf arg (=0)          exclude variants with MAF lower than this value
--project arg (=0)      project the new samples onto the existing PCs. Options are
                        0: disabled;
                        1: by multiplying the loadings with mean imputation for missing
    ↪ genotypes;
                        2: by solving the least squares system  $Vx=g$ . skip sites with
    ↪ missingness;
```

	3: by Augmentation, Decomposition and Procrusters ↪ transformation.
--inbreed arg (=0) ↪ structure. Options are	compute the inbreeding coefficient accounting for population
	0: disabled;
	1: compute per-site inbreeding coefficient and HWE test.
--ld-r2 arg (=0)	R2 cutoff for LD-based pruning (usually 0.2).
--ld-bp arg (=0) ↪ 1000000).	physical distance threshold in bases for LD window (usually
--ld-stats arg (=0)	statistics to compute LD R2 for pairwise SNPs. Options are
	0: the ancestry adjusted, i.e. correlation between residuals;
	1: the standard, i.e. correlation between two alleles.
--clump arg	assoc-like file with target variants and pvalues for clumping.
--clump-names arg (=CHR,BP,P) ↪ pvalue.	column names in assoc-like file for locating chr, pos and
--clump-p1 arg (=0.0001)	significance threshold for index SNPs.
--clump-p2 arg (=0.01)	secondary significance threshold for clumped SNPs.
--clump-r2 arg (=0.5)	r2 cutoff for LD-based clumping.
--clump-bp arg (=250000)	physical distance threshold in bases for clumping.

## 6.2 Which SVD method to use

This depends on your datasets, particularly the relationship between number of samples ( $N$ ) and the number of variants / features ( $M$ ) and the top PCs ( $k$ ). Here is an overview and the recommendation.

Method	Accuracy	Scenario
IRAM (-d 0)	Very high	$N < 1000$
Window-Based RSVD (-d 2)	Very high	$M > 1,000,000$
RSVD (-d 1)	High	accuracy insensitive
Full SVD (-d 3)	Exact	cost insensitive

## 6.3 Input formats

PCAone is designed to be extensible to accept many different formats. Currently, PCAone can work with SNP major genetic formats to study population structure. such as **PLINK**, **BGEN** and **Beagle**. Also, PCAone supports a comma delimited CSV format compressed by zstd, which is useful for other datasets requiring specific normalization such as single cell RNAs data.

## 6.4 Output formats

### 6.4.1 Eigen vectors

Eigen vectors are saved in file with suffix **.eigvecs**. Each row represents a sample and each col represents a PC.

### 6.4.2 Eigen values

Eigen values are saved in file with suffix **.eigvals**. Each row represents the eigenvalue of corresponding PC.

### 6.4.3 Features loadings

Features Loadings are saved in file with suffix **.loadings**. Each row represents a feature and each column represents a corresponding PC. Use **--printv** option to output it.

### 6.4.4 Variants information

A plink-like bim file named with **.mbim** is used to store the variants list with extra information. Currently, the **mbim** file has 7 columns with the 7th being the allele frequency. And PCAone only outputs this file whenever it's necessary to downstream analyses.

### 6.4.5 LD matrix

The matrix for calculating the ancestry-adjusted LD is saved in a file with suffix **.residuals**, and its associated variants information is stored in **mbim** file. For the binary file, the first 4-bytes stores the number of variants/SNPs, and the second 4-bytes stores the number of samples in the matrix. Then, the rest of the file is a sequence of **M** blocks of **N x 4** bytes each, where **M** is the number of variants and **N** is the number of samples. The first block corresponds to the first marker in the **.mbim** file, etc.

### 6.4.6 LD R2

The LD R2 for pairwise SNPs within a window can be outputted to a file with suffix **ld.gz** via **--print-r2** option. This file uses the same long format as the one **plink** used.



## 6.5 Memory-efficient modes

PCAone has both **in-core** and **out-of-core** mode for 3 different partial SVD algorithms, which are IRAM (`--svd 0`), Yu+Halko RSVD (`--svd 1`) and PCAone window-based RSVD (`--svd 2`). Also, PCAone supports full SVD (`--svd 3`) but with only **in-core** mode. Therefore, there are **7** ways for doing PCA in PCAone. In default PCAone uses **in-core** mode with `--memory 0`, which is the fastest way to do calculation. However, in case the server runs out of memory with **in-core** mode, the user can trigger **out-of-core** mode by specifying the amount of memory using `--memory` option with a value greater than 0.

### 6.5.1 Run PCAone window-based RSVD method (default) with in-core mode

```
./PCAone --bfile example/plink
```

### 6.5.2 Run PCAone window-based RSVD method (default) with out-of-core mode

```
./PCAone --bfile example/plink -m 2
```

### 6.5.3 Run Yu+Halko RSVD method with in-core mode

```
./PCAone --bfile example/plink --svd 1
```

### 6.5.4 Run Yu+Halko RSVD method with out-of-core mode

```
./PCAone --bfile example/plink --svd 1 -m 2
```

### 6.5.5 Run IRAM method with in-core mode

```
./PCAone --bfile example/plink --svd 0 -m 2
```

### 6.5.6 Run IRAM method with out-of-core mode

```
./PCAone --bfile example/plink --svd 0 -m 2
```

### 6.5.7 Run Full SVD method with in-core mode

```
./PCAone --bfile example/plink --svd 3
```

## 6.6 Data Normalization

PCAone will automatically apply the standard normalization for genetic data. Additionally, there are 3 different normalization method implemented with `--scale` option.

- 0: do just centering by subtracting the mean
- 1: do log transformation (usually for count data, such as bulk RNA-seq data)
- 2: do count per median log transformation (usually for single cell RNA-seq data)

One should choose proper normalization method for specific type of data.

## 6.7 Projection

Project new samples onto existing PCs is supported with `--project` option. First, we run PCAone on a set of reference samples and output the loadings:

```
PCAone -b ref_samples -k 10 --printv -o ref
```

Then, we need to read in the SNPs loadings from the ref set (`--read-V`) and its scaling factors (`--read-S`) as well as the allele frequencies from the `.mbim` file via `--match-bim`. **Note:** one can use the `--USV` option instead to simplify the usage since v0.4.8 Here is the example command to project new target samples and get the PCs coordinates.

```
PCAone -b new_samples \  
  --USV ref \  ## prefix to .eigvecs, .eigvals, .loadings  
  --project 2 \  ## check the manual on projection methods  
  -o new
```

## 6.8 HWE accounting for population structure

To test Hardy-Weinberg equilibrium in presence of population structure, we need to work on the so-called individual allele frequencies matrix  $\Pi$ , which can be reconstructed via the output of PCAone, i.e the `.eigvecs`, `.eigvals`, `.loadings` and `.mbim` files, generated by

```
PCAone -b example/plink -k 3 -V -o pcaone
```

Then we apply `--inbreed 1` option to obtain the P value of HWE and inbreeding coefficient per-site. The output file is named with suffix `.hwe`.

```
PCAone -b example/plink \  
  --USV pcaone \  
  --inbreed 1 \  
  -o inbreed
```

## 6.9 Ancestry-Adjusted LD matrix

LD patterns vary across diverse ancestry and structured groups, and conventional LD statistics, e.g. the implementation in `plink --ld`, failed to model the LD in admixed populations. Thus, we can use the so-called ancestry-adjusted LD statistics to account for population structure in LD. See our [paper](#) for more details.

To calculate the ancestry-adjusted LD matrix, we first figure out the number of principal components (`-k/--pc`) that capture population structure. In this example, assuming that 3 PCs can account for population structure, we enable `--ld` option to calculate and output the ancestry adjusted LD matrix in a file with suffix `.residuals`.

```
./PCAone -b example/plink -k 3 --ld -o adj
```

## 6.10 Report LD statistics

Currently, the LD R2 for pairwise SNPs within a window can be outputted via `--print-r2` option.

```
./PCAone -B adj.residuals \  
  --match-bim adj.mbam \  
  --ld-bp 1000000 \  
  --print-r2 \  
  -o adj
```

We provide the `calc_decay_bin.R` script to parse the output file `.ld.gz` and calculate the average R2 for each bin as well as plotting. We also provide the nextflow `ld.nf` for benchmarking the LD statistics.

## 6.11 Pruning based on Ancestry-Adjusted LD

Given the LD binary file `.residuals` and its associated variant file `.mbim`, we can do pruning based on user-defined thresholds and windows

```
./PCAone -B adj.residuals \  
  --match-bim adj.mbim \  
  --ld-r2 0.8 \  
  --ld-bp 1000000 \  
  -o adj
```

## 6.12 Clumping based on Ancestry-Adjusted LD

Likewise, we can do clumping based on the Ancestry-Adjusted LD matrix and user-defined association results

```
./PCAone -B adj_ld.residuals \  
  --match-bim adj.mbim \  
  --clump example/plink.pheno0.assoc,example/plink.pheno1.assoc \  
  --clump-p1 0.05 \  
  --clump-p2 0.01 \  
  --clump-r2 0.1 \  
  --clump-bp 10000000 \  
  -o adj
```

# 7 More tutorials

Let's download the example data first if you haven't done so.

```
wget http://popgen.dk/zilong/datahub/pca/example.tar.gz  
tar -xzf example.tar.gz && rm -f example.tar.gz
```

## 7.1 Genotype data (PLINK)

We want to compute the top 40 PCs for this genotype dataset using 20 threads and only 2 GBs memory. We will use the proposed window-based RSVD algorithm with default setting `--svd 2`.

```
./PCAone --bfile example/plink -k 40 -m 2 -n 20
```

Then, we can make a PCA plot in R.

```
pcs <- read.table("pcaone.eigvecs2",h=F)  
plot(pcs[,1:2+2], col=factor(pcs[,1]), xlab = "PC1", ylab = "PC2")  
legend("topright", legend=unique(pcs[,1]), col=1:4, pch = 21, cex=1.2)
```

## 7.2 Genotype dosage (BGEN)

Imputation tools usually generate the genotype probabilities or dosages in BGEN format. To do PCA with the imputed genotype probabilities, we can work on BGEN file with `--bgen` option instead.

```
./PCAone --bgen example/test.bgen -k 10 -n 4 -m 2
```

Then, we can make a PCA plot in R.

```
pcs <- read.table("pcaone.eigvecs",h=F)  
pop <- read.table("example/plink.fam",h=F)[,1]  
plot(pcs[,1:2], col=factor(pop), xlab = "PC1", ylab = "PC2")  
legend("topright", legend=unique(pop), col=factor(unique(pop)), pch = 21, cex=1.2)
```

### 7.3 Single cell RNA-seq data (CSV)

In this example, we run PCA for the single cell RNAs-seq data using a different input format with a normalization method called count per median log transformation (CPMED).

```
./PCAone --csv example/BrainSpinalCord.csv.zst -k 10 -n 20 -m 4 --scale 2 --svd 1
```

It should take around 5 minutes.

## 8 Acknowledgements

PCAone use [Eigen](#) for linear algebra operation. The IRAM method is based on [yixuan/spectra](#). The bgen lib is ported from [jeremycrae/bgen](#). The EMU and PCAngsd algorithms are modified from [@Jonas](#) packages.